

2016312761 여혁수 Homework 3 report

1-(a)

	Name	AlbumId	...	Bytes	UnitPrice
TrackId			...		
610	My Funny Valentine (Live)	49	...	29416781	0.99
620	Space Truckin'	50	...	39267613	0.99
621	Going Down / Highway Star	50	...	29846063	0.99
1581	Dazed And Confused	127	...	36052247	0.99
1666	Dazed And Confused	137	...	52490554	0.99
2429	We've Got To Get Together/Jingo	198	...	34618222	0.99
2432	Funky Piano	198	...	30200730	0.99
2819	Battlestar Galactica: The Story So Far	226	...	490750393	1.99
2820	Occupation / Precipice	227	...	1054423946	1.99
2821	Exodus, Pt. 1	227	...	475079441	1.99

1-(b)

Puja Srivastava

1-(c)

Led Zeppelin 14

Metallica 10

Deep Purple 11

Iron Maiden 21

Ozzy Osbourne 6

U2 10

1-(d)

The World of Classical Favourites 6

English Renaissance 6

2-(a)

	ArtistId	...	PlaylistId
0	1	...	1.0
1	1	...	8.0
2	1	...	17.0
3	1	...	1.0
4	1	...	8.0
...
8781	249	...	1.0
8782	249	...	5.0
8783	249	...	8.0
8784	249	...	12.0
8785	249	...	14.0

[8786 rows x 14 columns]

Columns of dataframe: ['ArtistId', 'ArtistName', 'AlbumId', 'Title', 'TrackId', 'TrackName',

'MediaTypeId', 'GenreId', 'Composer', 'Milliseconds', 'Bytes',

'UnitPrice', 'GenreName', 'PlaylistId']

2-(b)

['AC/DC', 'Accept', 'Led Zeppelin', 'Queen', 'Kiss', 'Deep Purple', 'Santana', 'Creedence Clearwater Revival', 'Foo Fighters', 'Guns N' Roses', 'Iron Maiden', 'Nirvana', 'Ozzy Osbourne', 'Pearl Jam', 'Red

Hot Chili Peppers', 'Skank', 'The Cult', 'The Rolling Stones', 'U2', 'Van Halen', 'Spyro Gyra', 'Miles Davis', 'Antônio Carlos Jobim', 'Caetano Veloso', 'Chico Science & Nação Zumbi', 'Various Artists', 'Gilberto Gil', 'Milton Nascimento', 'Cássia Eller', 'Djavan', 'Legião Urbana', 'Lulu Santos', 'Os Paralamas Do Sucesso', 'Tim Maia', 'Black Label Society', 'Black Sabbath', 'Metallica', 'Audioslave', 'Green Day', 'Faith No More', 'R.E.M.', 'Smashing Pumpkins', 'The Tea Party', 'Titãs', 'Eric Clapton', 'The Black Crowes', 'Cidade Negra', 'Jamiroquai', 'Amy Winehouse', 'Battlestar Galactica', 'Lost', 'The Office', 'English Concert & Trevor Pinnock', 'Eugene Ormandy', 'Michael Tilson Thomas & San Francisco Symphony', 'Berliner Philharmoniker & Herbert Von Karajan']

2-(c)

Top 7 genres:

	GenreId	COUNT(*)
0	1	1297
1	7	579
2	3	374
3	4	332
4	2	130
5	19	93
6	6	81

Construct a set of ten features for each artist:


	Rock ...	Number of playlists
AC/DC	18 ...	3
Accept	4 ...	4
Led Zeppelin	114 ...	3
Queen	45 ...	3
Kiss	35 ...	3
Deep Purple	92 ...	3

Santana	27 ...	3
Creedence Clearwater Revival	40 ...	2
Foo Fighters	33 ...	3
Guns N' Roses	28 ...	3
Iron Maiden	81 ...	4
Nirvana	29 ...	4
Ozzy Osbourne	18 ...	4
Pearl Jam	54 ...	4
Red Hot Chili Peppers	31 ...	3
Skank	23 ...	3
The Cult	30 ...	3
The Rolling Stones	41 ...	3
U2	112 ...	3
Van Halen	52 ...	3
Spyro Gyra	0 ...	3
Miles Davis	0 ...	3
Antônio Carlos Jobim	0 ...	3
Caetano Veloso	0 ...	4
Chico Science & Nação Zumbi	0 ...	3
Various Artists	0 ...	2
Gilberto Gil	0 ...	4
Milton Nascimento	0 ...	3
Cássia Eller	0 ...	3
Djavan	0 ...	4
Legião Urbana	0 ...	3
Lulu Santos	0 ...	2

Os Paralamas Do Sucesso	0 ...	3
Tim Maia	0 ...	3
Black Label Society	0 ...	2
Black Sabbath	0 ...	3
Metallica	0 ...	4
Audioslave	14 ...	3
Green Day	0 ...	3
Faith No More	15 ...	3
R.E.M.	14 ...	3
Smashing Pumpkins	0 ...	3
The Tea Party	0 ...	3
Titãs	0 ...	3
Eric Clapton	0 ...	4
The Black Crowes	0 ...	2
Cidade Negra	0 ...	3
Jamiroquai	10 ...	3
Amy Winehouse	0 ...	2
Battlestar Galactica	0 ...	2
Lost	0 ...	2
The Office	0 ...	2
English Concert & Trevor Pinnock	0 ...	6
Eugene Ormandy	0 ...	7
Michael Tilson Thomas & San Francisco Symphony	0 ...	5
Berliner Philharmoniker & Herbert Von Karajan	0 ...	6

[56 rows x 10 columns]

2-(d)



As you see the plot, the degree of cohesion goes higher. Based on 8, you can see that the slope of inertia score changes much gentler. It can be determined that the point 8 is the elbow point. So I think appropriate value of k is 8.

3

```
File Input Format Counters
  Bytes Read=467
File Output Format Counters
  Bytes Written=241
root@f4fd5fe71def:/opt/hadoop-2.7.1# hdfs dfs -cat output/*
1      dfsadmin
1      dfs.webhdfs.enabled
1      dfs.permissions.enabled
1      dfs.namenode.servicerpc
1      dfs.namenode.rpc
1      dfs.namenode.name.dir
1      dfs.namenode.https
1      dfs.namenode.http
1      dfs.datanode.use.datanode.hostname
1      dfs.client.use.datanode.hostname
root@f4fd5fe71def:/opt/hadoop-2.7.1#
```

Code)

```
import sqlite3
import pandas as pd
import matplotlib.pyplot as plt

conn = sqlite3.connect("Chinook_Sqlite.sqlite")
cursor = conn.cursor()

#1-(a)
print("1-(a)")
data = pd.read_sql_query("SELECT * FROM Track WHERE Milliseconds > 900000;", conn, index_col="TrackId")
print(data[:10])
print()

#1-(b)
print("1-(b)")
data = pd.read_sql_query("SELECT Invoice.CustomerId FROM Invoice INNER JOIN InvoiceLine ON Invoice.InvoiceId = InvoiceLine.InvoiceId GROUP BY CustomerId ORDER BY COUNT(InvoiceLine.TrackId);", conn, index_col="CustomerId")
min_amount = list(data.index)[0]
data = pd.read_sql_query("SELECT FirstName, LastName FROM Customer WHERE CustomerId='%d';" % min_amount, conn)
print(data['FirstName'][0], data['LastName'][0])
print()

#1-(c)
print("1-(c)")
data = pd.read_sql_query("SELECT COUNT(*), * FROM Album INNER JOIN Artist ON Album.ArtistId = Artist.ArtistId GROUP BY Album.ArtistId HAVING COUNT(*) > 5;", conn)
count_list = list(data['COUNT(*)'])
name_list = list(data['Name'])
for i in range(len(count_list)):
    print(str(name_list[i]) + ' ' + str(count_list[i]))
print()

#1-(d)
print("1-(d)")
data = pd.read_sql_query("SELECT Track.AlbumId, COUNT(DISTINCT PlaylistTrack.PlaylistId) FROM Track INNER JOIN PlaylistTrack ON Track.TrackId = PlaylistTrack.TrackId GROUP BY Track.AlbumId HAVING COUNT(DISTINCT PlaylistTrack.PlaylistId) > 5;", conn)
id_tuple = tuple(data['AlbumId'])
data2 = pd.read_sql_query("SELECT Album.Title FROM Album WHERE Album.AlbumId IN {}".format(id_tuple), conn)
for i in range(len(data)):
    print(data2['Title'][i], data['COUNT(DISTINCT PlaylistTrack.PlaylistId)'][i])
print()

#2-(a)
print("2-(a)")
df = pd.read_sql_query("SELECT * FROM Artist", conn)
df_album = pd.read_sql_query("SELECT * FROM Album", conn)
df = df.merge(df_album, how="outer", left_on="ArtistId", right_on="ArtistId")
df_track = pd.read_sql_query("SELECT * FROM Track", conn)
```

```

df = df.merge(df_track, how="outer", left_on="AlbumId", right_on="AlbumId")
df.rename(columns = {'Name_x' : 'ArtistName'}, inplace = True)
df.rename(columns = {'Name_y' : 'TrackName'}, inplace = True)
df_genre = pd.read_sql_query("SELECT * FROM Genre", conn)
df = df.merge(df_genre, how="outer", left_on="GenreId", right_on="GenreId")
df.rename(columns = {'Name' : 'GenreName'}, inplace = True)
df_playlistTrack = pd.read_sql_query("SELECT * FROM PlaylistTrack", conn)
df = df.merge(df_playlistTrack, how="outer", left_on="TrackId",
right_on="TrackId")
print(df)
print(df.columns)
print()

#2-(b)
print("2-(b)")
artist_dic = {}
artist_res = []
artist_res2 = []
for d in df.iterrows():
    artist_name = str(d[1][1])
    artist_id = d[1][0]
    if artist_name not in artist_dic:
        artist_dic[artist_name] = d[1][2]
    elif artist_name in artist_dic and artist_dic[artist_name] != d[1][2]:
        if artist_name not in artist_res:
            artist_res.append(artist_name)
            artist_res2.append(artist_id)

print(artist_res)
print()

#2-(c)
print("2-(c)")
data = pd.read_sql_query("SELECT GenreId, COUNT(*) FROM Track GROUP BY
GenreId ORDER BY COUNT(*) DESC", conn)
print("Top 7 genres: ")
print(data[:7])
id_tuple = tuple(data[:7]['GenreId'])
data = pd.read_sql_query("SELECT Genre.Name FROM Genre WHERE Genre.GenreId
IN {}".format(id_tuple), conn)
col_list = list(data['Name'])
genre_list = col_list
col_list.append("Number of albums")
col_list.append("Number of tracks")
col_list.append("Number of playlists")

new_df = pd.DataFrame(columns=col_list, index=artist_res)
for i in range(len(artist_res2)):
    for j in range(len(id_tuple)):
        data = pd.read_sql_query("SELECT COUNT(*) FROM Track INNER JOIN
Album ON Track.AlbumId = Album.AlbumId WHERE GenreId = {} AND
Album.ArtistId = {}".format(id_tuple[j], artist_res2[i]), conn)
        new_df[genre_list[j]][i] = list(data['COUNT(*)'])[0]

        data = pd.read_sql_query("SELECT COUNT(AlbumId) FROM Album WHERE
ArtistId = %d" % artist_res2[i], conn)
        new_df['Number of albums'][i] = list(data['COUNT(AlbumId)'])[0]
        data = pd.read_sql_query("SELECT COUNT(TrackId) FROM Track INNER JOIN
Album ON Track.AlbumId = Album.AlbumId WHERE Album.ArtistId = %d" %
artist_res2[i], conn)

```



```

    new_df['Number of tracks'][i] = list(data['COUNT(TrackId)'])[0]
    data = pd.read_sql_query("SELECT COUNT(DISTINCT PlaylistId) FROM Track
INNER JOIN Album ON Track.AlbumId = Album.AlbumId INNER JOIN PlaylistTrack
ON Track.TrackId = PlaylistTrack.TrackId WHERE Album.ArtistId = %d" %
artist_res2[i], conn)
    new_df['Number of playlists'][i] = list(data['COUNT(DISTINCT
PlaylistId)'])[0]

print(new_df)
print()

#2-(d)
print("2-(d)")
from sklearn.cluster import KMeans
ks = [2, 4, 6, 8, 10]
inertias = []
for k in ks:
    kmeans = KMeans(n_clusters=k)
    kmeans.fit(new_df)
    inertias.append(kmeans.inertia_)

plt.plot(ks, inertias, '-o')
plt.xlabel('number of clusters, k')
plt.ylabel('inertia score')
plt.xticks(ks)
plt.show()

conn.close()

```