

Week#13 Buffer Management in SQLite

Hyuksoo Yeo

2016312761

1. INTRODUCTION

Run TPC-C benchmark by varying the cache size. For example, change cache size to 50, 100, 150, 200. Then observe how TPS (txn/s) changes. Record and analyze the TPS for each transaction (DELIVERY, NEW_ORDER, ORDER_STATUS, PAYMENT, STOCK_LEVEL). Lastly, present and analyze experimental results.

2. METHODS

For this experiment, install SQLite library and python. Because we will use pytpcc benchmark. Then setup pytpcc benchmark by preparing the SQLite configuration file and manipulate the path to the SQLite database. Then load TPC-C database and run to evaluating TPC-C benchmark. Repeating running by changing page cache size. We can change the size by editing the value in --buffer in command.

3. Performance Evaluation

3.1 Experimental Setup

System setup:

Type	Specification
OS	Ubuntu 20.04.3 LTS
CPU	Intel® Core™ i3-9100F CPU @ 3.60GHz
Memory	16GB
Kernel	5.11.0-27-generic
Data Device	Western Digital WD Blue 500GB
Log Device	Western Digital WD Blue 500GB

Benchmark setup:

Type	Configuration
DB size	1GB (10 warehouse)
Buffer Pool Size	300MB (30% of DB size)
Benchmark Tool	tpcc-mysql
Runtime	1200s
Connections	8

3.2 Experimental Results

Cache size = 50 (pages)

```
yhs@yhs-VirtualBox:~/SWE3033-F2021/week-13/pytpcc$ python tpcc.py --warehouse=10
--config=./sqlite.config --no-load --duration=1800 --buffer=50 sqlite
3.31.1
journal mode delete
cache_size 50
11-28-2021 19:20:20 [<module>:240] INFO : Initializing TPC-C benchmark using Sql
iteDriver
11-28-2021 19:20:20 [execute:056] INFO : Executing benchmark for 1800 seconds
=====
Execution Results after 1800 seconds
-----

```

	Executed	Time (µs)	Rate
DELIVERY	4875	203413285.494	23.97 txn/s
NEW_ORDER	53806	1035843201.16	51.94 txn/s
ORDER_STATUS	4921	8661961.07864	568.12 txn/s
PAYMENT	51595	527723878.145	97.77 txn/s
STOCK_LEVEL	4873	13237733.1257	368.11 txn/s
TOTAL	120070	1788880059.0	67.12 txn/s

```
-----
```

Cache size = 100

```
yhs@yhs-VirtualBox:~/SWE3033-F2021/week-13/pytpcc$ python tpcc.py --warehouse=10
--config=./sqlite.config --no-load --duration=1800 --buffer=100 sqlite
3.31.1
journal mode delete
cache_size 100
11-28-2021 19:56:06 [<module>:240] INFO : Initializing TPC-C benchmark using Sql
iteDriver
11-28-2021 19:56:06 [execute:056] INFO : Executing benchmark for 1800 seconds
=====
Execution Results after 1800 seconds
-----

```

	Executed	Time (µs)	Rate
DELIVERY	5055	168905270.1	29.93 txn/s
NEW_ORDER	56287	1066466877.7	52.78 txn/s
ORDER_STATUS	4982	7875113.72566	632.63 txn/s
PAYMENT	53622	532592760.801	100.68 txn/s
STOCK_LEVEL	5045	12825190.5441	393.37 txn/s
TOTAL	124991	1788665212.87	69.88 txn/s

```
-----
```

Cache size = 150

```
yhs@yhs-VirtualBox:~/SWE3033-F2021/week-13/pytpcc$ python tpcc.py --warehouse=10
--config=./sqlite.config --no-load --duration=1800 --buffer=150 sqlite
3.31.1
journal mode delete
cache_size 150
11-28-2021 20:43:27 [<module>:240] INFO : Initializing TPC-C benchmark using Sql
iteDriver
11-28-2021 20:43:27 [execute:056] INFO : Executing benchmark for 1800 seconds
=====
Execution Results after 1800 seconds
-----

```

	Executed	Time (µs)	Rate
DELIVERY	4803	169011243.343	28.42 txn/s
NEW_ORDER	54590	1070619347.81	50.99 txn/s
ORDER_STATUS	4876	6813793.42079	715.61 txn/s
PAYMENT	52162	532573827.744	97.94 txn/s
STOCK_LEVEL	4933	9299520.73097	530.46 txn/s
TOTAL	121364	1788317733.05	67.86 txn/s

```
-----
```

Cache size = 200

```
yhs@yhs-VirtualBox:~/SWE3033-F2021/week-13/pytpcc$ python tpcc.py --warehouse=10
--config=./sqlite.config --no-load --duration=1800 --buffer=200 sqlite
3.31.1
journal mode delete
cache_size 200
11-28-2021 21:18:06 [<module>:240] INFO : Initializing TPC-C benchmark using Sql
iteDriver
11-28-2021 21:18:06 [execute:056] INFO : Executing benchmark for 1800 seconds
=====
Execution Results after 1800 seconds
-----

```

	Executed	Time (µs)	Rate
DELIVERY	4743	166993699.551	28.40 txn/s
NEW_ORDER	52298	1085202093.36	48.19 txn/s
ORDER_STATUS	4642	5009377.00272	926.66 txn/s
PAYMENT	50097	525443260.431	95.34 txn/s
STOCK_LEVEL	4666	4809072.49451	970.25 txn/s

TOTAL	116446	1787457502.84	65.15 txn/s

	cache size (pages)	50	100	150	200
transaction					
DELIVERY		23.97	29.93	28.42	28.40
NEW_ORDER		51.94	52.78	50.99	48.19
ORDER_STATUS		568.12	632.63	715.61	926.66
PAYMENT		97.77	100.68	97.94	95.34
STOCK_LEVEL		368.11	393.37	530.46	970.25
TOTAL		67.12	69.88	67.86	65.15

In case of DELIVERY, NEW_ORDER, and PAYMENT, TPS is highest when cache size is 100 pages and decreases by increasing cache size after 100 pages. While the case of ORDER_STATUS and STOCK_LEVEL, TPS continuously goes up by increasing cache size.

TOTAL TPS is highest at cache size = 100.

4. Conclusion

In this experiment, I learned how to manage buffer in SQLite by using pytpcc benchmark. TPS is different at each transaction type and cache size.

5. REFERENCES

[1] <https://github.com/meeeeejin/SWE3033-F2021/tree/main/week-13>